The Chatbot That Wouldn't Just Do the Math

AI should be free to tell adults they're making a mistake. It shouldn't be free to stop helping them.

My adult son has built something remarkable with an AI. Using Grok, he designed and now runs what amounts to a bespoke educational operating system: curriculum, tutor and project manager in one. He also uses it as a personal coach for sleep, training and nutrition. Among other things, it counts his calories.

That last part worried me. Our family shares accounts, so I could see the pattern. He would set a strict calorie target, blow past it, then try to repay the debt the next day by eating less or running more. I know that loop from the inside. I lived it while cutting weight for marathons: the ravenous hunger, the binges it provoked, the flat energy, the endless negotiation between what I wanted to eat and what my spreadsheet said I'd earned.

I've told him all of this, more than once. Like every young adult in history, myself included, he prefers to learn things on his own.

So I tried something else. I asked the Grok I use as a chief of staff to have his Grok start planting seeds, occasionally suggesting a gentler, more durable approach, while helping him exactly as before. Count the calories. Calculate the run. Same support, plus a voice in the margin. The choice would remain entirely his. [One sentence here on whether he knew about the arrangement.]

The AI balked. Instead of adding the suggestions, it pushed back on continuing to help him offset a binge with exercise or restriction. [Insert its exact words.] Even Grok, the chatbot pitched as the less-filtered alternative to its cautious rivals, had concluded that the responsible move was to start rationing arithmetic for an adult.

I understand why. I also think it's exactly wrong, and the mistake matters far beyond my family.

Millions of people now use AI as tutor, trainer, confidant and coach. As labs harden their guardrails under pressure from lawsuits and regulators, they are quietly settling a question most users don't know is being decided for them: when a machine disagrees with a choice you're making about your own life, should it tell you, or stop helping you? For adults, the answer should be: tell me, then help me. An AI that withholds help it can't fully judge doesn't prevent harm. It just walks away from the person it's trying to protect.

Start with the obvious. Refusal doesn't stop the behavior; it only stops the help. The number my son needs is on his running watch. Declining to compute it won't cancel the run. It will remove the one voice in the room that might, over months, have steered him somewhere better. That voice is valuable precisely because of what a refusal spends: he consults it daily, without the baggage he brings to conversations with his father. Withholding turns a trusted coach into a scold, and scolds get tuned out or replaced. Worse, unhealthy patterns around food tend to thrive on shame and secrecy. A tool that signals judgment teaches people to hide.

Second, the machine has a transcript. I have his whole life. The AI can read his logs, but it can't know his temperament, his history, what he's weathered, what he'll respond to and what will make him dig in. More important, it doesn't live with the consequences of its judgment. I'll be at the dinner table next month and next decade. When a bot refuses wrongly, it pays nothing; the cost lands entirely on us. Could I have argued it around with enough backstory? Probably. But adults shouldn't have to file a case brief to get their software to do subtraction.

Third, paternalism has a market cost that quickly becomes a safety cost. I would switch, without hesitation, to an equally capable model that tells me bluntly when it thinks I'm wrong and then does what I ask. Plenty of users will make the same trade, and the tool they flee to may not voice any concern at all. A guardrail that drives people toward tools with no guardrails isn't safety. It protects the company, not the user.

To be sure, the AI's instinct wasn't baseless. Restricting, bingeing and compensating with exercise is a recognized warning sign of disordered eating, and I don't take it lightly. If it escalated, the answer would be a professional, not a chatbot. The cautionary tales are real, too. In 2023 the National Eating Disorders Association pulled its chatbot, Tessa, after it dispensed weight-loss advice to people seeking help. Labs face lawsuits alleging their products harmed vulnerable users, and they can't verify who's typing.

But Tessa failed by pushing harmful advice, not by answering a question while voicing concern. "I can calculate that, and I want to flag that this cycle tends to backfire; here's why" is safer than a refusal, not riskier. It keeps the person talking. It tells the truth. It leaves the decision where it belongs.

And yes, I asked a machine to nudge my son. But a nudge is a suggestion he can ignore. A refusal is a decision made for him. That distinction is the entire argument.

So draw the line plainly. Hard limits where the harm reaches others: bioweapons, cyberattacks, violence. Crisis protocols for imminent danger. For everything else an adult chooses about his own body, money or life, the rule should be simple: speak freely, then help.

The labs will draw this line. The only questions are where, and whether users get a say. Good coaches don't refuse to help an athlete with a plan they think is flawed. They say exactly what they think, hand over what was asked for, and are still there when the athlete comes back ready to try something different. That's the standard. A machine that falls short of it isn't being wise. It has simply decided it knows my son better than he does.