Step 1: Thesis Extraction

Core thesis: When adults make choices about their own lives, AI should be free to disagree but never to withhold help, because a machine that refuses help it can't fully judge doesn't prevent harm. It just walks away from the person it's trying to protect.

What I cut, and why:
- **The case against calorie counting.** It's contestable, since self-monitoring is a core component of evidence-based behavioral weight programs, and it's beside the point. The essay is about AI refusal, not nutrition. If you argue diet science, you'll lose readers who track calories happily.
- **"Any and all attempts to align AI... will almost certainly backfire spectacularly."** You walk this back two sentences later. It's an overclaim that invites easy dismissal.
- **The extended sympathy for AI labs.** Compressed into the "To Be Sure" paragraph, where it does work instead of diluting your argument.
- **"I don't want to explain every request to AI."** As written, it reads as entitlement. I reframed it: adults shouldn't have to file a case brief to get arithmetic.

Step 2: Pressure-Test

1. The bot may have been right, and readers will say so first. This is your biggest vulnerability. Restricting, bingeing, and compensating with exercise is a textbook warning pattern. "Excessive exercise" is listed as a compensatory behavior in the DSM-5 criteria for bulimia nervosa. A skeptical reader's first reaction: "The machine spotted a possible eating disorder and the dad got mad." There's also a damaging precedent. In 2023, the National Eating Disorders Association pulled its chatbot, Tessa, after it gave weight-loss advice to people seeking help.
- **Fortify:** Don't argue the concern was wrong. Argue the response was wrong. Concede the pattern is worrying; your marathon experience gives you standing to say so. Then show that refusal is the wrong intervention. The number is on his watch, withholding signals judgment, and judgment breeds secrecy. Distinguish Tessa: it failed by pushing harmful advice, not by answering questions while voicing concern. Signal that you'd involve a professional if this escalated. That tells readers you're clear-eyed, and you should mean it.

2. The hypocrisy trap. You're writing an essay about autonomy built on an anecdote in which you read your adult son's conversations and asked an AI to steer him, apparently without his knowledge. A hostile reader will argue the bot was protecting his autonomy from you.
- **Fortify:** Confront it directly. Your request was suggestions he could ignore, not restrictions. That nudge-versus-veto distinction is your thesis, so use it. Also state plainly whether he knows the account is shared and knew about the arrangement. If he didn't, the essay won't survive contact with readers.
- **Non-negotiable:** Get his consent before publishing. This is his health behavior, and he's identifiable through your byline.

3. "AI lacks context" is only half right. The bot arguably had more context than usual: your account plus his logs. Context also cuts both ways. If the bot can't trust its own judgment, it can't verify yours either. It has no way to know you're a caring parent rather than someone manipulating a vulnerable person.
- **Fortify:** The real asymmetry is accountability, not information. You'll be at the dinner table for decades; the bot faces no consequences for a wrong refusal. Add the asymmetry of errors: wrongly helping with arithmetic costs little, while wrongly refusing costs the relationship.

4. Your line is mushy. "Err on the side of generous autonomy," plus "perhaps" even in egregious cases, gets you dismissed as either an absolutist or a waffler.
- **Fortify:** Commit to a line:
- Hard limits where the harm reaches others.
- Crisis protocols for imminent danger.
- For self-regarding adult choices, voice concern but don't withhold.

5. The evidence is one family and one bot.
- **Fortify:** Quote the bot's pushback verbatim. That's the most persuasive sentence you can add. Also use the irony that Grok is marketed as the less-filtered chatbot. If even Grok drifts toward paternalism, this is a default across the industry, not a glitch.

6. The labs' constraints are real. They face lawsuits, vulnerable users, and no way to verify who's typing.
- **Fortify:** Acknowledge this, then argue that "concern without withholding" is the cheaper and safer design, not the riskier one.

Step 3: Op-Ed Structure

- **Lede:** The son's remarkable AI-built education system, then the calorie loop you recognize from your marathon days, then the "plant seeds" request, then the bot balks. The Grok irony lands the hook.
- **Nut graph:** Millions use AI as coach and confidant, and labs are hardening guardrails right now. The question being settled is whether a machine that disagrees should tell you or stop helping you. Answer: tell me, then help me.
- **Evidence:**
1. Refusal stops the help, not the behavior, and it spends the trust that made the bot a good messenger.
2. The bot has a transcript; you have his whole life, plus accountability.
3. Paternalism drives users to tools with no guardrails at all.
- **To Be Sure:** The warning sign is real, Tessa happened, and labs face lawsuits. But candor beats refusal. Then own the nudge: a nudge is a suggestion, a refusal is a decision made for him. Draw the line.
- **Kicker:** Good coaches state their concerns, hand over what was asked, and stay. Anything less isn't wisdom.

Step 4: The Draft

The Chatbot That Wouldn't Just Do the Math

AI should be free to tell adults they're making a mistake. It shouldn't be free to stop helping them.

My adult son has built something remarkable with an AI. Using Grok, he designed and now runs what amounts to a bespoke educational operating system: curriculum, tutor and project manager in one. He also uses it as a personal coach for sleep, training and nutrition. Among other things, it counts his calories.

That last part worried me. Our family shares accounts, so I could see the pattern. He would set a strict calorie target, blow past it, then try to repay the debt the next day by eating less or running more. I know that loop from the inside. I lived it while cutting weight for marathons: the ravenous hunger, the binges it provoked, the flat energy, the endless negotiation between what I wanted to eat and what my spreadsheet said I'd earned.

I've told him all of this, more than once. Like every young adult in history, myself included, he prefers to learn things on his own.

So I tried something else. I asked the Grok I use as a chief of staff to have his Grok start planting seeds, occasionally suggesting a gentler, more durable approach, while helping him exactly as before. Count the calories. Calculate the run. Same support, plus a voice in the margin. The choice would remain entirely his. [One sentence here on whether he knew about the arrangement.]

The AI balked. Instead of adding the suggestions, it pushed back on continuing to help him offset a binge with exercise or restriction. [Insert its exact words.] Even Grok, the chatbot pitched as the less-filtered alternative to its cautious rivals, had concluded that the responsible move was to start rationing arithmetic for an adult.

I understand why. I also think it's exactly wrong, and the mistake matters far beyond my family.

Millions of people now use AI as tutor, trainer, confidant and coach. As labs harden their guardrails under pressure from lawsuits and regulators, they are quietly settling a question most users don't know is being decided for them: when a machine disagrees with a choice you're making about your own life, should it tell you, or stop helping you? For adults, the answer should be: tell me, then help me. An AI that withholds help it can't fully judge doesn't prevent harm. It just walks away from the person it's trying to protect.

Start with the obvious. Refusal doesn't stop the behavior; it only stops the help. The number my son needs is on his running watch. Declining to compute it won't cancel the run. It will remove the one voice in the room that might, over months, have steered him somewhere better. That voice is valuable precisely because of what a refusal spends: he consults it daily, without the baggage he brings to conversations with his father. Withholding turns a trusted coach into a scold, and scolds get tuned out or replaced. Worse, unhealthy patterns around food tend to thrive on shame and secrecy. A tool that signals judgment teaches people to hide.

Second, the machine has a transcript. I have his whole life. The AI can read his logs, but it can't know his temperament, his history, what he's weathered, what he'll respond to and what will make him dig in. More important, it doesn't live with the consequences of its judgment. I'll be at the dinner table next month and next decade. When a bot refuses wrongly, it pays nothing; the cost lands entirely on us. Could I have argued it around with enough backstory? Probably. But adults shouldn't have to file a case brief to get their software to do subtraction.

Third, paternalism has a market cost that quickly becomes a safety cost. I would switch, without hesitation, to an equally capable model that tells me bluntly when it thinks I'm wrong and then does what I ask. Plenty of users will make the same trade, and the tool they flee to may not voice any concern at all. A guardrail that drives people toward tools with no guardrails isn't safety. It protects the company, not the user.

To be sure, the AI's instinct wasn't baseless. Restricting, bingeing and compensating with exercise is a recognized warning sign of disordered eating, and I don't take it lightly. If it escalated, the answer would be a professional, not a chatbot. The cautionary tales are real, too. In 2023 the National Eating Disorders Association pulled its chatbot, Tessa, after it dispensed weight-loss advice to people seeking help. Labs face lawsuits alleging their products harmed vulnerable users, and they can't verify who's typing.

But Tessa failed by pushing harmful advice, not by answering a question while voicing concern. "I can calculate that, and I want to flag that this cycle tends to backfire; here's why" is safer than a refusal, not riskier. It keeps the person talking. It tells the truth. It leaves the decision where it belongs.

And yes, I asked a machine to nudge my son. But a nudge is a suggestion he can ignore. A refusal is a decision made for him. That distinction is the entire argument.

So draw the line plainly. Hard limits where the harm reaches others: bioweapons, cyberattacks, violence. Crisis protocols for imminent danger. For everything else an adult chooses about his own body, money or life, the rule should be simple: speak freely, then help.

The labs will draw this line. The only questions are where, and whether users get a say. Good coaches don't refuse to help an athlete with a plan they think is flawed. They say exactly what they think, hand over what was asked for, and are still there when the athlete comes back ready to try something different. That's the standard. A machine that falls short of it isn't being wise. It has simply decided it knows my son better than he does.

(~1,000 words, leaving room for the two bracketed insertions.)

Before You File

1. **Get your son's sign-off.** This is non-negotiable, for the essay's credibility and for your relationship.
2. **Pull the verbatim pushback** from the Grok logs. One real quote beats any paragraph I can write.

3. **Fill in the transparency bracket honestly.** If he didn't know about the arrangement, tell him before you publish, not after.
4. **Fact-check the Tessa details and the Grok marketing language** against current sources before submission.