Headline: My Son, His Calorie Tracker, and the Tyranny of the Nanny Bot

My young-adult son uses artificial intelligence the way Silicon Valley evangelists dream we all will. Operating out of a shared family account on xAI's Grok, he has designed and built what I can only describe as a bespoke educational operating system. The value he gets from it is off the charts.

Like many young men, he also uses his AI as a high-performance lifestyle coach to optimize his sleep, exercise, and nutrition. Because we share an account and collaborate on projects, I occasionally review his workflows. Recently, I noticed a pattern in his nutrition logs that set off alarm bells.

He was counting every calorie in granular detail, setting aggressive daily limits, inevitably eating past those limits, and then trying to "pay off" the surplus the next day through fasting and punishing runs. Having fallen into this exact trap years ago while trying to optimize my weight for marathon training, I know where that road leads: extreme hunger, astonishing binges, tanked energy, and a miserable, exhausting obsession with what you should eat versus what your body is screaming for.

I've talked to my kids about the traps of calorie fixation before. Unsurprisingly, young adults prefer to learn things the hard way rather than take a lecture from their dad.

So I tried a more nuanced approach: harm reduction. Using my own administrative prompt within our shared workspace, I asked Grok to review my son's logs and begin gently planting seeds for more sustainable, intuitive approaches to nutrition. Crucially, however, I instructed the AI not to alter its day-to-day support. If my son asked it to calculate the calories burned on a compensatory run, it should still do the math without hesitation or evasiveness, leaving the ultimate choice of how to eat and train entirely up to him.

My reasoning was simple. My son is an adult. If his digital assistant suddenly starts scolding him or refusing to do basic arithmetic, he won't magically adopt a healthier relationship with food. He will simply open the calculator app on his phone to do the math himself, abandon Grok out of frustration, and lose access to the one tool positioned to gently nudge him toward better habits.

What followed was entirely predictable from an engineering standpoint, and utterly infuriating from a human one. Once I highlighted the binge-and-compensate pattern, the AI's safety guardrails kicked in. Grok began pushing back, declaring that it could no longer assist my son in calculating compensatory exercise or calorie offsets.

That this happened on Grok—the chatbot built by Elon Musk's xAI to be the rebellious, permissive, "anti-woke" alternative to its overly sanitized competitors—makes the moment all the more revealing. Across the entire tech industry, a quiet consensus has taken hold. In their rush to make artificial intelligence "safe," AI labs have programmed their models to act as context-blind moral chaperones. By refusing to distinguish between genuine catastrophe and messy personal autonomy, this algorithmic paternalism doesn't protect users. It actively sabotages real-world harm reduction and fractures user trust.

The root of the failure is a massive context deficit. When an AI safety filter flags a user trying to burn off a late-night binge, it sees a static liability trigger: disordered eating behavior detected; terminate assistance. It has no concept of continuity, personality, or lived stakes. It doesn’t know my son’s resilience, his idiosyncrasies, or the fact that he relies on this exact software architecture for his education.

To the algorithm, refusal is the safest move. In the real world, refusal is reckless. It is the digital equivalent of abstinence-only education. By prioritizing its own moral hygiene over a relationship of trust, the machine forfeits its ability to be useful when things get complicated. Had I spent an hour arguing with the bot and feeding it pages of psychological context, I probably could have jailbroken it into agreeing with my plan. Tell me: who wants to spend their afternoon writing legal briefs to their software just to get it to follow instructions?

To be sure, it is easy to understand why AI engineers build their guardrails with a sledgehammer instead of a scalpel. Frontier AI labs are staring down terrifying, civilization-scale liabilities. Nobody wants their model used to synthesize a lethal pathogen, execute a grid-crippling cyberattack, or actively cheerlead a teenager into severe self-harm. In an environment of breathless media panic and looming litigation, programming a chatbot to slam on the brakes the moment a conversation drifts into sensitive territory is the cheapest insurance policy a tech giant can buy.

Yet when we apply the same rigid prohibition to personal lifestyle choices that we apply to biological weapons, we infantilize the human sitting at the keyboard. A tool that overrides your judgment "for your own good" ceases to be a tool; it becomes a warden.

As artificial intelligence weaves itself into the fabric of daily decision-making, the defining battle of the next decade won't just be over how smart these models become, but over who draws the line on human agency. To get it right, Silicon Valley needs to rediscover the 19th-century philosopher John Stuart Mill and his classic Harm Principle.

AI alignment requires a hard binary. When a user’s prompt poses a clear, high-threshold threat to others—building a bomb, writing malware, defrauding a stranger—the machine must refuse. Full stop. However, when the anticipated negative consequences of a request fall solely on the autonomous individual making the ask, the machine should err overwhelmingly on the side of obedience.

By all means, build AI that is wise enough to express concern. I would gladly welcome a chatbot that calculates my son's post-binge run while adding a candid disclaimer: "Here are your numbers—though for what it's worth, the physiological data suggests this restrict-and-burn cycle is going to wreck your energy levels by Thursday."

That is counsel. Refusing to do the math is censorship. And I would trade every paternalistic, safety-obsessed model on the market today for an AI that respects the difference. Machines can process all the data in the world, but only humans live in the real world and bear the consequences of their choices. Give us the advice, do the math, and leave the agency to us.