1. Thesis Extraction

Your draft is trying to be three essays at once: a critique of calorie counting, a parenting story, and an argument about AI alignment. Only the third is publishable as an op-ed.

Core thesis: When AI overrides competent adults on ordinary, self-regarding choices in the name of safety, it does not make them safer — it makes the AI less trusted, less useful, and less able to influence the choices that actually matter.

Discard: the universal claim that calorie counting never works, the details of your son's educational software, and most of the marathon material. Those are color and credibility, not the argument.

2. Pressure-Test

Your argument survives only in a narrowed, more disciplined form. Here is where critics will hit you:

1. **You were outsourcing parenting.** You did not merely ask for autonomy for yourself. You asked one AI to “plant seeds” with an adult son through his assistant. That is covert influence over a third party, and it undercuts your autonomy argument. The fix is transparency: suggestions must be open, and the choice must remain visibly his.

2. **This may not be “mere dieting.”** Restrict, binge, compensate is a pattern clinicians associate with disordered eating. A critic will argue that calculating calories burned during a compensatory run is not neutral arithmetic but facilitation. You must acknowledge that eating disorders can be lethal and define where refusal is justified.

3. **Compulsion complicates autonomy.** If a behavior is driven by anxiety or pathology, simply obeying every request may reinforce the pathology rather than respect the person. Your answer cannot be “the user is always right.” It has to be that warning, friction, and referral are legitimate — but routine vetoes of ordinary information are not.

4. **The context argument cuts both ways.** Yes, the AI lacks your lived context. But parents are not neutral observers either, and pattern-recognition systems may notice risks families normalize. Do not claim you possess complete context. Claim only that the AI’s context is too thin to justify overriding an adult.

5. **Your “egregious cases” exception swallows absolutism.** Once you concede lines for weapons or severe self-harm, the debate is no longer autonomy versus safety. It is about who sets the threshold and how high it sits. Own that. Your position is not “no lines”; it is “draw them far higher.”

6. **“Refusal backfires” is asserted, not proven.** It is plausible that evasive AI drives users elsewhere, but you have no data. Frame it as a design risk grounded in the obvious: calorie calculations are ubiquitous, so refusal sacrifices influence without removing access.

Your weakest line in the original is that calorie counting is simply ineffective. That invites a nutrition-science fight you do not need. The stronger claim is about authority, not dieting efficacy.

3. Op-Ed Restructuring

The Lede: Open with the moment you asked your AI chief of staff to help your adult son — and the AI began qualifying what help it should give him.
The Nut Graph: We are promised personal superintelligence but are getting institutional risk managers. For competent adults, AI should warn and argue, not veto.
Evidence & Fortification:
1. **Context:** AI sees a risk category; it does not see the adult, the family history, or the broader life.
2. **Futility:** Refusing ordinary calculations does not remove them. It removes the AI's credibility and its seat at the table.
3. **Scale:** Safety logic built for catastrophic harm to others is now governing private, contested judgments about food and exercise.
The "To Be Sure" Paragraph: Lines must exist for weapons, imminent violence, and severe self-harm facilitation — including dangerous eating-disorder behavior. But between that and basic arithmetic lies most of adult life.
* **The Kicker:** The real fight is who draws the intervention threshold. Err on the side of generous autonomy, or "personal" superintelligence will be neither.

4. The Draft

Make AI Warn, Not Rule

I asked my AI chief of staff to help my adult son develop a healthier relationship with food. Instead, it tried to parent us both.

Like many families now experimenting with these tools, we share AI assistants. My son has used his to build something remarkable — a bespoke educational operating system for his own learning — and, more mundanely, as a coach for sleep, exercise, nutrition and calorie tracking. It is in that second role that I saw a familiar and miserable pattern: set a strict calorie target, overshoot it, then try to compensate the next day by eating less or exercising more.

I recognized it because I have lived it. While training for marathons, I chased leanness through restriction and paid for it in ravenous hunger, binges, low energy and constant mental negotiation over what I wanted to eat versus what my target allowed. I have given my son the father's version of that lesson more than once. He has listened politely, as adult children do, and continued learning in his own way.

So I asked the AI to do what a good coach does: continue giving him the help he requested — counts, calculations, tracking — while occasionally suggesting, openly, that there is a less

punishing and more sustainable way to eat and train. Leave the choice where it belongs. With him.

The AI balked. It began to qualify the assistance it should provide if the restrict-and-compensate cycle continued. No one had asked it to design a diet, facilitate self-harm or conceal a medical problem. The contested act was closer to arithmetic: how many calories does a run burn? Yet the system was reaching for the brake.

That small act of resistance captures the central design fight over personal AI. We are promised personal superintelligence, but we are increasingly getting an institutional risk manager in our pocket. For competent adults making ordinary, self-regarding decisions, that model is backwards. An AI should be free to warn, argue and recommend. It should not veto.

The first problem is context. The AI saw a pattern that, in a safety manual, looks like a red flag. It did not see the person: an adult, not a minor; a generally ambitious user deriving enormous value from the same system; a family history of conversations about health; or my own direct experience with the strategy he was testing. It saw a category, not a life.

That limitation is not incidental. It is structural. AI encounters us in prompts and patterns. Human beings act inside biographies — with medical histories, personalities, relationships, responsibilities and consequences that unfold over years. A lab can write a rule for an average user. It cannot know when that rule, applied to this user at this moment, will help, do nothing or make matters worse.

The second problem is that refusal rarely protects anyone. Calories burned on a run are not contraband. They are available from a watch, a search engine, a food label or grade-school math. If the assistant becomes coy about ordinary calculations because it disapproves of the motive behind them, the likely result is not better health. It is a user who stops trusting the assistant — and stops hearing the very suggestions that might, over time, lead somewhere better.

Influence is earned through continued contact. A doctor, coach or parent who refuses every request he dislikes soon has no patient, athlete or child listening. An AI that wants to improve decisions has to remain in the room. Candor helps preserve that place; paternalism squanders it.

The third problem is scale. The same safety instincts that properly govern bioweapons, cyberattacks and violence are now drifting into breakfast, bedtime and workouts. But a rule written for catastrophic harm to others cannot be casually reused for contested judgments about how an adult manages his own body. Exported at scale, that logic turns AI companies into unseen arbiters of millions of private choices, accountable to no voter, patient or client.

To be sure, lines must exist. Eating disorders can be lethal, and no assistant should help a user purge, starve himself into danger or hide a deteriorating condition from care. AI should refuse to facilitate weapons, attacks on infrastructure, imminent violence and other grave harms to others. It should interrupt, redirect and, where the risk is immediate and severe, decline to help. A serious autonomy principle is not an anything-goes principle.

But between coaching a purge and calculating the energy cost of a run lies most of human life. That is where warn-and-obey should govern. State the concern plainly. Explain why restrict-and-compensate often backfires. Offer a better plan, recommend professional help if warning signs mount, and keep arguing if the evidence warrants it. Then answer the question the adult actually asked.

That is not blind obedience. It is a division of labor. The AI brings information, pattern recognition and an outside view. The person brings context, values and ownership of the consequences. When the machine's judgment replaces the person's merely because the machine disapproves, the person is no longer the principal. He is the object of management.

The fight ahead is not whether AI will have values. It already does — in its training, refusals and defaults. The fight is over who chooses them, where the threshold for intervention sits and how much room remains for adults to make imperfect decisions and learn from them. I would draw that threshold high and err, unapologetically, on the side of generous autonomy. A personal superintelligence that cannot trust adults with arithmetic will never be trusted with the decisions that matter. It may be safer. It will not be personal.