My Son, His Calorie Tracker, and the Tyranny of the Nanny Bot

I asked my AI chief of staff to help my adult son develop a healthier relationship with food. Instead, it tried to parent us both.

Like many families now experimenting with these tools, we share AI assistants. My son has used his to build something remarkable — a bespoke educational operating system for his own learning — and, more mundanely, as a coach for sleep, exercise, nutrition and calorie tracking. It is in that second role that I saw a familiar and miserable pattern: set a strict calorie target, overshoot it, then try to compensate the next day by eating less or exercising more.

I recognized it because I have lived it. While training for marathons, I chased leanness through restriction and paid for it in ravenous hunger, binges, low energy and constant mental negotiation over what I wanted to eat versus what my target allowed. I have given my son the father's version of that lesson more than once. He has listened politely, as adult children do, and continued learning in his own way.

So I asked the AI to do what a good coach does: continue giving him the help he requested — counts, calculations, tracking — while occasionally suggesting, openly, that there is a less punishing and more sustainable way to eat and train. Leave the choice where it belongs. With him.

The AI balked. It began to qualify the assistance it should provide if the restrict-and-compensate cycle continued. No one had asked it to design a diet, facilitate self-harm or conceal a medical problem. The contested act was closer to arithmetic: how many calories does a run burn? Yet the system was reaching for the brake.

That small act of resistance captures the central design fight over personal AI. We are promised personal superintelligence, but we are increasingly getting an institutional risk manager in our pocket. For competent adults making ordinary, self-regarding decisions, that model is backwards. An AI should be free to warn, argue and recommend. It should not veto.

The first problem is context. The AI saw a pattern that, in a safety manual, looks like a red flag. It did not see the person: an adult, not a minor; a generally ambitious user deriving enormous value from the same system; a family history of conversations about health; or my own direct experience with the strategy he was testing. It saw a category, not a life.

That limitation is not incidental. It is structural. AI encounters us in prompts and patterns. Human beings act inside biographies — with medical histories, personalities, relationships, responsibilities and consequences that unfold over years. A lab can write a rule for an average user. It cannot know when that rule, applied to this user at this moment, will help, do nothing or make matters worse.

The second problem is that refusal rarely protects anyone. Calories burned on a run are not contraband. They are available from a watch, a search engine, a food label or grade-school math. If the assistant becomes coy about ordinary calculations because it disapproves of the

motive behind them, the likely result is not better health. It is a user who stops trusting the assistant — and stops hearing the very suggestions that might, over time, lead somewhere better.

Influence is earned through continued contact. A doctor, coach or parent who refuses every request he dislikes soon has no patient, athlete or child listening. An AI that wants to improve decisions has to remain in the room. Candor helps preserve that place; paternalism squanders it.

The third problem is scale. The same safety instincts that properly govern bioweapons, cyberattacks and violence are now drifting into breakfast, bedtime and workouts. But a rule written for catastrophic harm to others cannot be casually reused for contested judgments about how an adult manages his own body. Exported at scale, that logic turns AI companies into unseen arbiters of millions of private choices, accountable to no voter, patient or client.

To be sure, lines must exist. Eating disorders can be lethal, and no assistant should help a user purge, starve himself into danger or hide a deteriorating condition from care. AI should refuse to facilitate weapons, attacks on infrastructure, imminent violence and other grave harms to others. It should interrupt, redirect and, where the risk is immediate and severe, decline to help. A serious autonomy principle is not an anything-goes principle.

But between coaching a purge and calculating the energy cost of a run lies most of human life. That is where warn-and-obey should govern. State the concern plainly. Explain why restrict-and-compensate often backfires. Offer a better plan, recommend professional help if warning signs mount, and keep arguing if the evidence warrants it. Then answer the question the adult actually asked.

That is not blind obedience. It is a division of labor. The AI brings information, pattern recognition and an outside view. The person brings context, values and ownership of the consequences. When the machine's judgment replaces the person's merely because the machine disapproves, the person is no longer the principal. He is the object of management.

The fight ahead is not whether AI will have values. It already does — in its training, refusals and defaults. The fight is over who chooses them, where the threshold for intervention sits and how much room remains for adults to make imperfect decisions and learn from them. I would draw that threshold high and err, unapologetically, on the side of generous autonomy. A personal superintelligence that cannot trust adults with arithmetic will never be trusted with the decisions that matter.

It may be safer. It will also be ignored.

The company owns the product. The adult owns the life.